Fooling and Protecting Deep learning models

Divyam Madaan

Twitter: @divyam3897
Github: @divyam3897

About me

- ➔ Final year undergraduate
- ➔ Deep Learning Researcher at FOR.ai
- ➔ Working on Adversarial Training and Robustness

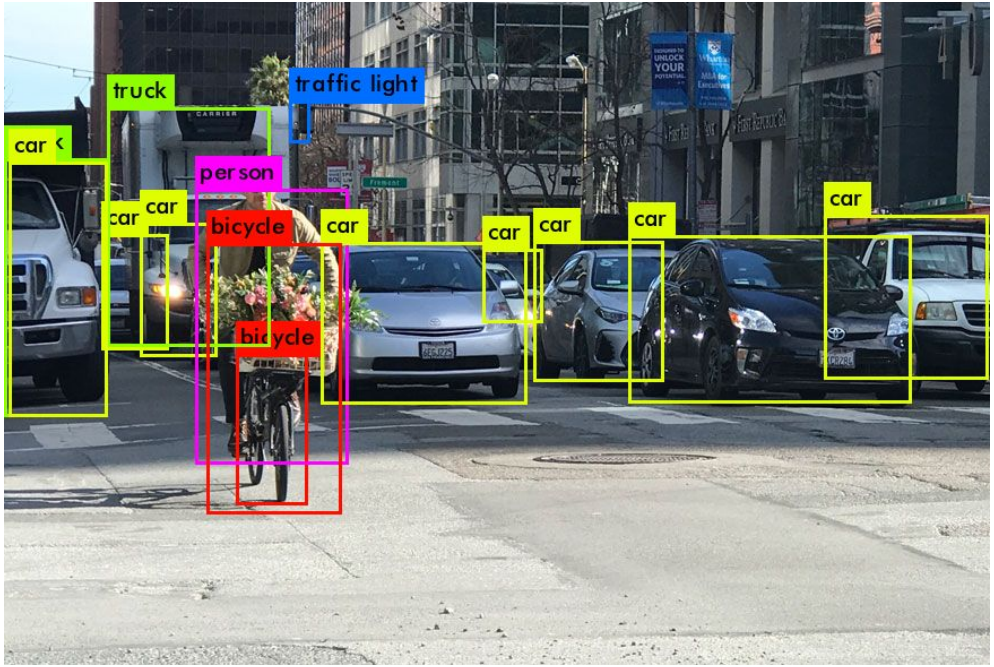

Machine learning reached state of the art on various tasks

Good models make mistakes

Attack on the machine learning pipeline

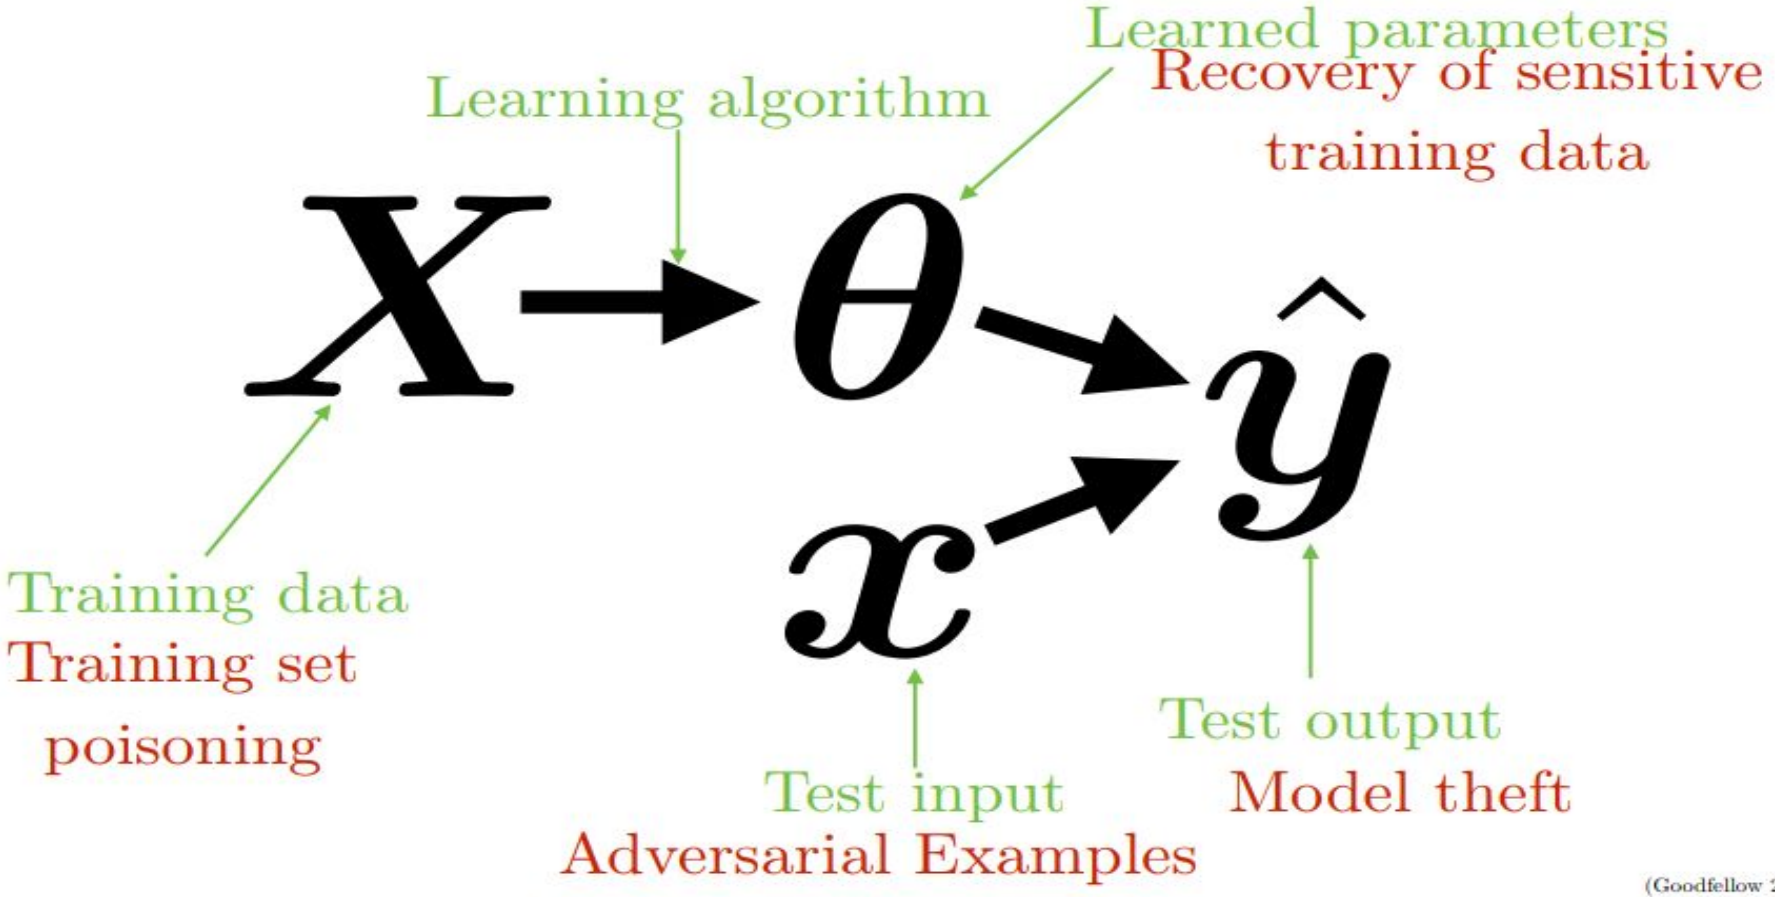

(Goodfellow 2018)

Timeline of Learning Security

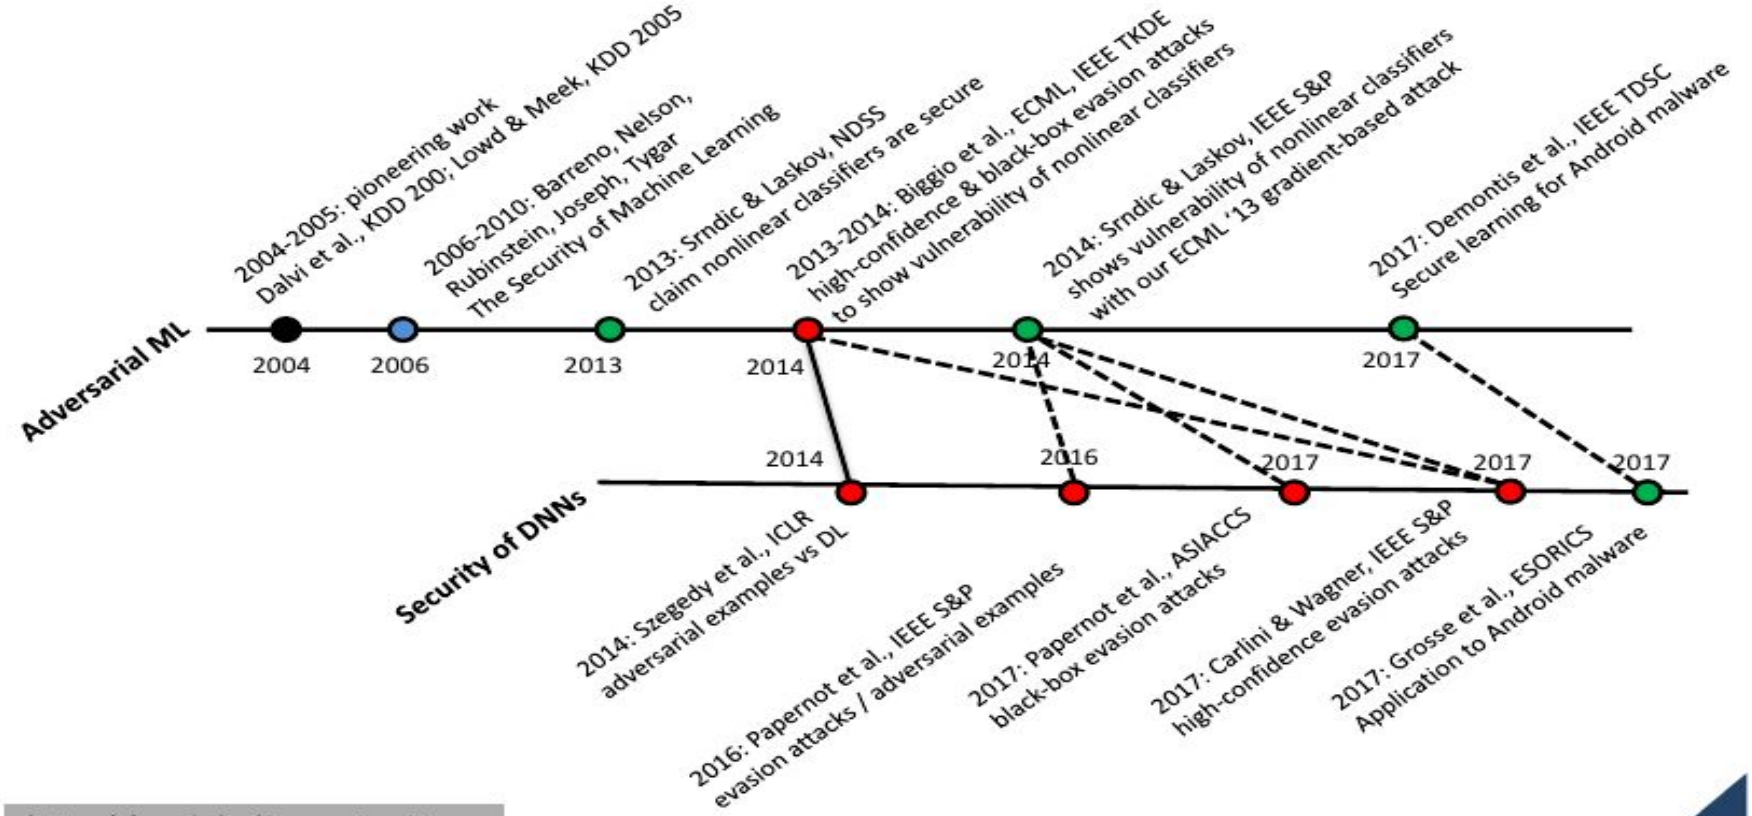

What this means for us?

- Deep learning algorithms (Machine learning in general) are susceptible to attacks
- Use with caution in critical deployments
- Evaluate a model's adversarial resilience - not just accuracy/precision/recall
- Spend effort to make model robust to tempering

Defending the machines

- Distillation (Train model 2x, feed first DNN output logits into second DNN input layer)
- Train models with adversarial samples i.e ironing out imperfect knowledge learnt in the model)
- Special regularization methods/loss functions (simulating adversarial content during trainiₙg)

Timeline of Defences

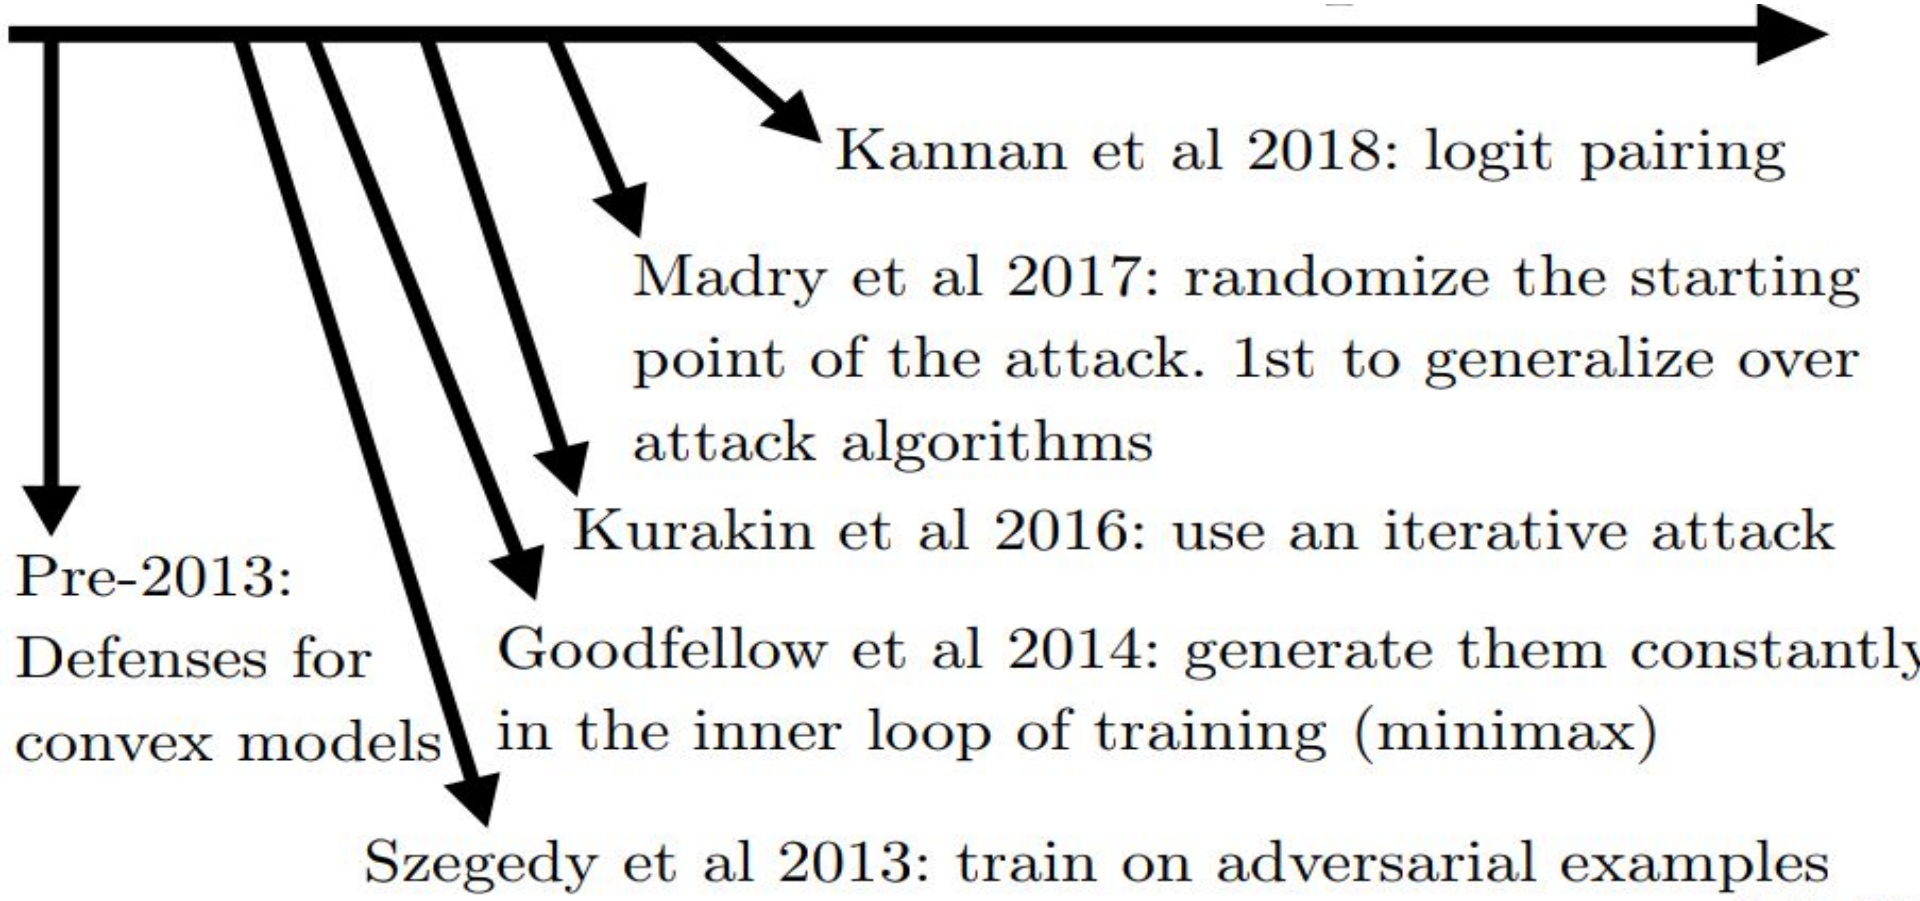

(Goodfellow 2018)


LIVE DEMO

Thank you
Questions?

Twitter: @divyam3897

Linkedin: https://www.linkedin.com/in/dmadaan3897

Github: @divyam3897